

# Code Execution with MCP: Concepts & Experimentation Guide

Created by Paul Brower on 2026-01-24 11:59 AM

Based on Anthropic Engineering: *Code Execution with MCP*

## Overview

Traditional AI tool usage loads all tool definitions upfront into the context window, consuming significant tokens. **Code execution with MCP** takes a different approach: instead of exposing tool definitions directly to the model, it presents MCP integrations as programmatic APIs that agents invoke through generated code.

**Key benefit:** Token usage can drop from ~150,000 tokens to ~2,000 tokens (98.7% reduction) by loading only necessary tool definitions on demand.

## Core Concepts

### 1. File-Based Tool Discovery

Tools are organized in a filesystem structure that agents navigate like a codebase:

```
servers/
  └── google-drive/
      ├── getDocument.ts
      └── searchFiles.ts
  └── salesforce/
      ├── updateRecord.ts
      └── query.ts
  └── slack/
      ├── postMessage.ts
      └── getChannelHistory.ts
```

Each tool is a TypeScript file with typed interfaces. Agents explore this structure to find relevant tools rather than having everything preloaded.

### 2. Tool Wrapper Pattern

Individual tools wrap MCP calls with type definitions:

```
export async function getDocument(
  input: GetDocumentInput
): Promise<GetDocumentResponse> {
  return callMCPTool<GetDocumentResponse>(
    'google_drive__get_document',
    input
  );
}
```

### 3. Agent-Generated Code

Instead of direct tool calls, agents write executable code that composes multiple tools:

```
// Fetch a document from Google Drive
const transcript = (await gdrive.getDocument({
  documentId: 'abc123'
})).content;

// Update a Salesforce record with the content
await salesforce.updateRecord({
  objectType: 'SalesMeeting',
  recordId: '00Q5f000001abcXYZ',
  data: { Notes: transcript }
});
```

## 4. In-Process Data Filtering

Large datasets are filtered within the execution environment, not passed through the model:

```
const allRows = await gdrive.getSheet({ sheetId: 'abc123' });
const pendingOrders = allRows.filter(row =>
  row["Status"] === 'pending'
);
console.log(`Found ${pendingOrders.length} pending orders`);
```

## 5. Control Flow in Code

Complex operations use standard programming constructs:

```
let found = false;
while (!found) {
  const messages = await slack.getChannelHistory({
    channel: 'C123456'
});
  found = messages.some(m => m.text.includes('deployment complete'));
  if (!found) await new Promise(r => setTimeout(r, 5000));
}
```

## 6. State Persistence

Agents maintain state across executions using filesystem operations:

```
// Save intermediate results
const leads = await salesforce.query({
  query: 'SELECT Id, Email FROM Lead LIMIT 1000'
});
await fs.writeFile('./workspace/leads.csv', csvData);

// Later execution can resume from saved state
const saved = await fs.readFile('./workspace/leads.csv');
```

# Security Considerations

## Sandboxed Execution

Agent-generated code must run in a secure environment with:

- Process isolation
- Resource limits (CPU, memory, network)
- Monitoring and logging

## Data Privacy via Tokenization

The MCP client can automatically tokenize PII before it reaches the model:

```
// Agent sees tokenized data
{ email: '[EMAIL_1]', phone: '[PHONE_1]', name: '[NAME_1]' }

// Real data flows between services untokenized
```

## Deterministic Security Rules

Define explicit data flow policies specifying which systems can access particular data types.

## Benefits Summary

Benefit	Description
<b>Token Efficiency</b>	Load only necessary tool definitions on demand
<b>Scalability</b>	Handle hundreds of tools across dozens of MCP servers
<b>Reduced Latency</b>	Loops and conditionals execute in code, not via repeated model calls
<b>Complex Data Ops</b>	Aggregations, joins, filtering happen in-process
<b>Composability</b>	Multiple MCP servers compose naturally through code

## Experimentation Steps

### Phase 1: Understand the Current MCP Setup

- **Step 1.1:** Review your existing MCP configuration

```
cat ~/.claude/settings.json
ls -la ~/.claude/memory/
```

- **Step 1.2:** Examine how your memory MCP servers work

- Look at the tool definitions they expose
  - Understand the current direct tool-call pattern
- **Step 1.3:** Document current token usage patterns
    - Note how many tools are loaded in typical sessions
    - Identify scenarios where tool count becomes unwieldy

### Phase 2: Create a Simple Tool Wrapper Library

- **Step 2.1:** Create a `servers/` directory structure in this project

```
mkdir -p servers/memory
```

- **Step 2.2:** Create TypeScript wrapper for memory server tools

```
// servers/memory/createEntities.ts
interface Entity {
```

```

    name: string;
    entityType: string;
    observations: string[];
}

export async function createEntities(entities: Entity[]): Promise<void> {
    // Wrapper that would call mcp_memory_create_entities
}

```

- **Step 2.3:** Create an index file that exports available tools

```

// servers/memory/index.ts
export * from './createEntities';
export * from './searchNodes';
export * from './readGraph';

```

### Phase 3: Build a Code Execution Environment

- **Step 3.1:** Set up a sandboxed TypeScript execution environment
  - Consider using Docker or a Node.js sandbox library
  - Implement resource limits
- **Step 3.2:** Create a simple REPL that can execute tool wrapper code

```

// executor.ts
import * as memory from './servers/memory';

async function executeAgentCode(code: string) {
    // Sandbox and execute the code with access to tool wrappers
}

```

- **Step 3.3:** Test basic code execution with your wrappers

### Phase 4: Implement Progressive Tool Discovery

- **Step 4.1:** Create a tool search/discovery interface

```

// discovery.ts
export async function searchTools(query: string): Promise<ToolDefinition[]> {
    // Search through servers/ directory for matching tools
}

```

- **Step 4.2:** Test discovering tools on-demand vs preloading all definitions
- **Step 4.3:** Measure token savings in a controlled experiment

### Phase 5: Claude Code Integration Exploration

- **Step 5.1:** Explore how Claude Code's existing MCP integration works
  - Review .mcp.json configuration format
  - Understand how tools are exposed to the model
- **Step 5.2:** Prototype a "code execution" MCP server
  - Single tool: execute\_code that runs TypeScript in sandbox
  - Access to tool wrapper library within sandbox
- **Step 5.3:** Test the pattern with a real workflow
  - Example: Query memory, filter results, update specific nodes

- Compare token usage vs direct tool calls

## Phase 6: Advanced Patterns

- **Step 6.1:** Implement state persistence
  - Create a workspace/ directory for intermediate results
  - Test resuming operations from saved state
- **Step 6.2:** Build reusable “skills”
  - Save working code snippets as functions
  - Create a skill library that grows over time
- **Step 6.3:** Explore PII tokenization
  - Implement a data sanitization layer
  - Test that sensitive data doesn’t reach the model

---

## Questions to Answer Through Experimentation

1. **Token Efficiency:** How much context can you save with on-demand tool loading?
  2. **Latency Trade-offs:** Does code execution add latency vs direct tool calls for simple operations?
  3. **Complexity Threshold:** At what number of tools does code execution become clearly beneficial?
  4. **Security Model:** What sandboxing approach works best for your use case?
  5. **Developer Experience:** How does this pattern feel compared to direct tool invocation in Claude Code?
  6. **Composability:** What multi-tool workflows become easier with code execution?
- 

## Relevant to Your Claude Code Usage

Based on your current setup with memory MCP servers, here are specific opportunities:

1. **Memory Operations:** Complex queries across your knowledge graph could be written as code rather than multiple tool calls
  2. **Batch Operations:** Creating/updating many entities could be a single code execution vs many individual calls
  3. **Conditional Logic:** “If entity exists, update it; otherwise create it” is natural in code
  4. **Data Transformation:** Filter and transform query results in-process before storing
- 

## Resources

- Model Context Protocol Documentation
- Anthropic Engineering Blog
- MCP Server Examples