

# Code Execution with MCP

Created by Paul Brower on 2026-01-24 11:59 AM

**Dramatically speed up AI-assisted batch operations by having the model write code instead of making individual tool calls.**

This repository explores the Code Execution with MCP pattern from Anthropic's engineering blog, with working examples you can run yourself.

## The Problem

When you ask an AI assistant to process records one at a time, it's painfully slow:

```
User: "Update the status on these 50 records"  
  
AI: [calls tool to find record 1]  
AI: [calls tool to update record 1]  
AI: [calls tool to find record 2]  
AI: [calls tool to update record 2]  
... repeat 48 more times ...
```

Each tool call requires a full round-trip: model reasoning → API call → tool execution → response parsing → model reasoning again. In our testing, **each tool call takes ~4 seconds**. For 50 records with 2 calls each, that's **~7 minutes**.

## The Solution

Instead of the AI calling tools one-by-one, have it write code that does the batch operation:

```
// One tool call that runs this code:  
for (const id of recordIds) {  
  const record = await memory.openNodes([id]);  
  await memory.addObservations([{ entityName: id, contents: ['status: processed'] }]);  
}
```

**Result: 50 records in 1 millisecond instead of 7 minutes.**

## Why This Matters

Benefit	Description
<b>Speed</b>	Eliminate model round-trips for iterative operations
<b>Token Efficiency</b>	Load fewer tool definitions into context ( 95% reduction)
<b>Reduced Context Rot</b>	Less back-and-forth means cleaner conversation history
<b>Natural Composition</b>	Loops, conditionals, and data transformations are natural in code
<b>Scalability</b>	Handle hundreds of tools without bloating context

## When to Use Each Approach

Use Direct Tool Calls For	Use Code Execution For
Exploration and discovery	Batch operations (“for each X, do Y”)
One-off queries	Complex filtering and transformation
Learning the data structure	Multi-step workflows
Simple operations	Operations on many records

## Experiments

### Experiment 1: Batch Record Updates

Tests the performance difference between direct tool calls and code execution for iterative memory operations.

#### Results Summary

Metric	Direct Tool Calls	Code Execution	Improvement
5 iterations	42.7 seconds	<1 ms	42,000x faster
50 iterations	7 minutes	1 ms	430,000x faster
Tool calls per batch	2 per iteration	1 total	90-99% reduction
Tool tokens in context	6,000	300	95% reduction

#### Key Finding

**Model round-trips are the bottleneck, not tool execution.** Each direct tool call takes ~4.3 seconds (model reasoning + API latency + MCP execution). Code execution eliminates this for iterative operations.

## Running the Experiments Yourself

### Prerequisites

- Node.js 18+
- Claude Code CLI
- npm

### Setup

#### 1. Clone the repository

```
git clone https://github.com/yourusername/code-mode.git  
cd code-mode
```

#### 2. Install dependencies

```
npm install
```

#### 3. Build the project

```
npm run build
```

## 4. Configure MCP servers

Create or update `.mcp.json` in the project root:

```
{  
  "mcpServers": {  
    "code_mode_memory": {  
      "type": "stdio",  
      "command": "npx",  
      "args": ["-y", "@modelcontextprotocol/server-memory"],  
      "env": {  
        "MEMORY_FILE_PATH": "/path/to/code-mode/code-mode-memory.jsonl"  
      }  
    },  
    "code_executor": {  
      "type": "stdio",  
      "command": "node",  
      "args": ["dist/mcp-server/index.js"],  
      "cwd": "/path/to/code-mode"  
    }  
  }  
}
```

Replace `/path/to/code-mode` with your actual project path.

## 5. Restart Claude Code to pick up the new MCP configuration

### Running Tests

Test the memory wrapper library:

```
npm test
```

Test the code executor:

```
npm run test:executor
```

### Using the Code Executor in Claude Code

Once configured, you can use the `execute_code` tool in Claude Code:

```
// Example: Count active records  
const graph = await memory.readGraph();  
const activeCount = graph.entities.filter(entity =>  
  entity.observations.some(obs => obs === "status: active")  
).length;  
console.log(`Active records: ${activeCount}`);
```

The code executes in a sandboxed environment with access to memory operations but restricted system access.

## Project Structure

```
code-mode/  
  └── servers/memory/      # TypeScript wrappers for memory MCP tools
```

```
|- types.ts          # Type definitions
|- client.ts        # MCP client interface
|- operations.ts    # Tool wrapper functions
|- index.ts         # Public exports
|-- executor/
|   |- sandbox.ts   # VM-based sandboxed execution
|   |- index.ts     # Executor entry point
|-- mcp-server/
|   |- index.ts     # MCP server exposing execute_code tool
|-- test/           # Test files
`-- code-mode-memory.jsonl # Test data (50 sample records)
Experiment1.md      # Detailed experiment documentation
mcp-config.example.json # Example MCP configuration
```

## How It Works

### 1. Tool Wrappers

Instead of exposing raw MCP tools, we create typed TypeScript wrappers:

```
// Direct MCP call (verbose, untyped)
mcp_code_mode_memory_read_graph()

// Wrapper (clean, typed)
await memory.readGraph()
```

### 2. Sandboxed Execution

Agent-generated code runs in a Node.js VM sandbox with:

- Access to memory tool wrappers
- Console output capture
- Timeout protection
- No filesystem or network access

### 3. Single Tool Call

The `execute_code` MCP tool accepts JavaScript code and runs it in the sandbox:

```
Model → execute_code("...loop over 50 records...") → Sandbox executes all
iterations → Single result
```

vs. direct tool calls:

```
Model → tool call → result → Model → tool call → result → ... (100 round trips)
```

## Resources

- Anthropic Engineering: Code Execution with MCP
- Model Context Protocol Documentation
- MCP Server Examples

# **License**

MIT